

Technische Universität Berlin

Notes

BIUB

Sascha Lange

Meta Blub

Let \mathcal{S} denote state space and \mathcal{A} denote action space. Let $\mathcal{I} = \{1, \dots, P\}$ be the set of players and $\mathcal{P} = \{N_1, \dots, N_P\}$ the set of corresponding (neural network) function approximators, where $N_i : \mathcal{S} \mapsto \mathcal{A}$, $N_i(s) = a$, corresponds to player i .

Playing with a new teammate p , the goal is to predict the action a that p will take, facing a new state s . Based on the pool \mathcal{P} of pretrained agents, we will define the notion of *consistency* of two states, that will help us building a simple classifier to do so.

Definition 1 (consistency). *For $s, s' \in \mathcal{S}$, let*

$$\text{consistency}_{\mathcal{P}}(s, s') = \frac{1}{P} \cdot |\{i \in \mathcal{P} : N_i(s) = N_i(s')\}|.$$

An agent $i \in \mathcal{I}$ is said to be consistent in (s, s') , if $N_i(s) = N_i(s')$.

We now define a simple distance function, which we can minimize later, to predict an action that p will take:

Definition 2 (Distance). *For $s, s' \in \mathcal{S}$, let the distance be defined as*

$$d(s, s') = 1 - \text{consistency}_{\mathcal{P}}(s, s').$$

After having observed data $D = \{(s_1, a_1), \dots, (s_n, a_n)\}$ from a new teammate p , we can give the action classifier as

$$C_p(s) = \underset{a}{\operatorname{argmin}} \{d(\tilde{s}, s) : (\tilde{s}, a) \in D\}.$$

So we guess as the action p takes in s , the action p took in the state $s_{i \leq n}$, with maximum $\text{consistency}_{\mathcal{P}}(s_{i \leq n}, s)$. However the suggested distance, does not account for the fact, that some players in \mathcal{I} are more likely similar¹ to p , than others. When the consistency voting happens, we would like to give those agents more weight than others. Therefore, we introduce the following notion on the observed data D :

Definition 3 (agreement). *For $s, s' \in \mathcal{S}, i, p \in \mathcal{P}$, let*

$$\text{agreement}_D(i, p) = \frac{2}{n(n-1)} \cdot |\{(s, s') \in D \times D : N_i(s) = N_i(s') \text{ and } N_p(s) = N_p(s')\}|.$$

The agreement of two agents i, p is higher, the more pairs $(s, s') \in D$ exist, that i and p are both consistent on. Under the assumption, that this implies the agents are more likely to perform the same actions in new states, we now define p-similarity of two states, as a weighted version of consistency:

Definition 4 (p-similarity). *For $s, s' \in \mathcal{S}, i, p \in \mathcal{P}$, let*

$$\text{similarity}_p(s, s') = \frac{1}{P} \sum \{\text{agreement}_D(i, p) \in \mathbb{R} : i \text{ is consistent in } (s, s')\}. \quad (1)$$

Let C_θ denote our meta classifier. We want to learn $\theta = \theta_0$, such that

¹similar w.r.t to which states they consider consistent

- after small number L of gradient steps on data D from agent A , to obtain θ_L , the network C_{θ_L} performs well on predicting actions of A

So we obtain updated network params after $i \leq L$ steps on D from A by

$$\theta_i^A = \theta_{i-1}^A - \alpha \Delta_{\theta} \mathcal{L}_A(C_{\theta_{i-1}^A})$$

for a **single** Task A , and thus the meta-objective becomes

$$\sum_{A \in \text{POOL}} \mathcal{L}_P(C_{\theta_L}^A) =: \mathcal{L}_{\text{Meta}},$$

where \mathcal{L}_P denotes the loss on the hold out set corresponding to A . Both A and P are agents, but A denotes agents at training time and P denotes agents at test time, indicating that players can be humans. Note however, that **the different notation simply denotes disjoint data, but from the same agent $P=A$.**

Finally we have the outer loop update given by

$$\theta_0 = \theta_0 - \beta \Delta \mathcal{L}_{\text{Meta}}.$$

Using the idea of incorporating implicit soft cluster assignment (see slack) into the learning process we may obtain for C_{θ} the following architecture:

